

Week 4 Assignment: AI in Software Engineering

Theme: *"Building Intelligent Software Solutions"*

- **Q1:** Explain how AI-driven code generation tools (e.g., GitHub Copilot) reduce development time. What are their limitations?

AI-driven code generation tools such as GitHub Copilot accelerate software development by providing intelligent code suggestions based on context, comments, and previous code patterns. They reduce development time by automating boilerplate code writing, offering quick implementations of known algorithms, and reducing context-switching for developers. This allows engineers to focus on architecture and logic rather than syntax and repetitive tasks.

However, limitations exist. First, Copilot may produce insecure or incorrect code due to lack of understanding of business logic or context. It sometimes "hallucinates" APIs that don't exist or misuses functions. Second, its suggestions can introduce subtle bugs if blindly accepted. Lastly, ethical concerns arise when generated code resembles open-source code without proper attribution.

Developers must critically evaluate Copilot's output and use it as a productivity aid, not a replacement for software design or review processes.

- **Q2:** Compare supervised and unsupervised learning in the context of automated bug detection.

Supervised learning involves labelled data — for bug detection, this means datasets where code samples are explicitly marked as buggy or clean. With such datasets, classifiers (e.g., Random Forests, SVMs) can be trained to predict bugs in unseen code based on features like complexity, syntax, or commit history.

Unsupervised learning, by contrast, does not require labels. It's useful for detecting anomalies in large codebases or log files — for example, clustering similar code patterns and flagging outliers that deviate from the norm. Techniques like K-Means or autoencoders help identify unusual patterns or error logs that could indicate bugs.

In summary, supervised learning is precise and effective when labelled data is available, while unsupervised learning is useful for discovering hidden bugs in unlabelled, dynamic environments.

- **Q3:** Why is bias mitigation critical when using AI for user experience personalization?

Bias mitigation is essential in AI-driven UX personalization to ensure fair, inclusive, and ethical digital experiences. Personalization systems learn from user data, which may reflect societal or platform-specific biases — such as under-representing certain groups, reinforcing stereotypes, or favouring majority preferences.

If left unchecked, biased AI could result in unequal visibility (e.g., job recommendations skewed toward one gender) or discriminatory interactions. Bias can originate from imbalanced training datasets, model assumptions, or unmonitored feedback loops.

Therefore, applying fairness techniques such as re-weighting training data, introducing fairness constraints during model training, and continuously auditing outputs helps ensure all users receive equitable treatment. Tools like IBM AI Fairness 360 or Google's What-If Tool support this process.

Case Study Analysis

How does AIOps improve software deployment efficiency? Provide two examples.

AIOps (Artificial Intelligence for IT Operations) combines machine learning, analytics, and automation to enhance software deployment and operations. It helps DevOps teams manage complex infrastructure by providing real-time insights, predictive failure analysis, and automated remediation.

Example 1: Predictive Incident Management

AIOps tools analyse log and performance data to detect anomalies before they escalate into outages. For instance, memory leak patterns identified during staging can trigger preventive alerts or automated scaling before reaching production.

Example 2: Self-Healing Pipelines

AIOps platforms monitor CI/CD pipelines and automatically resolve known errors. If a deployment script fails due to a common syntax error, the system can rollback the pipeline and apply a fix based on prior incidents.

These capabilities reduce downtime, increase deployment speed, and free developers to focus on feature delivery rather than firefighting infrastructure issues.